

# Package ‘bigmemory’

July 2, 2014

**Version** 4.4.6

**Date** 2013-11-18

**Title** Manage massive matrices with shared memory and memory-mapped files

**Author** Michael J. Kane <kaneplusplus@gmail.com>, John W. Emerson <jayemerson@gmail.com>, and Peter Haverty <haverty.peter@gene.com>

**Maintainer** Michael J. Kane <bigmemoryauthors@gmail.com>

**Contact** Mike and Jay <bigmemoryauthors@gmail.com>

**Depends** methods, utils, bigmemory.sri, BH

**Enhances** biganalytics, bigtabulate, synchronicity

**LinkingTo** BH

**Suggests** RUnit

**Description** Create, store, access, and manipulate massive matrices.  
Matrices are allocated to shared memory and may use memory-mapped files. Packages biganalytics, bigtabulate, synchronicity, and bigalgebra provide advanced functionality.

**OS\_type** unix

**License** LGPL-3

**Copyright** (C) 2013 John W. Emerson, Michael J. Kane, and Peter Haverty

**URL** <http://www.bigmemory.org>

**LazyLoad** yes

**NeedsCompilation** yes

**Repository** CRAN

**Date/Publication** 2013-11-18 15:47:41

R topics documented:

|  |           |
|--|-----------|
| bigmemory-package . . . . .                            | 2         |
| as.big.matrix-methods . . . . .                        | 4         |
| big.matrix, filebacked.big.matrix, etc... . . . .      | 5         |
| big.matrix-class . . . . .                             | 9         |
| big.matrix.descriptor-class . . . . .                  | 11        |
| deepcopy . . . . .                                     | 12        |
| describe, attach.big.matrix, attach.resource . . . . . | 13        |
| flush . . . . .  | 14        |
| morder, mpermute . . . . .                             | 15        |
| mwhich . . . . .                                       | 17        |
| mwhich-methods . . . . .                               | 18        |
| sub.big.matrix, is.sub.big.matrix . . . . .            | 19        |
| write.big.matrix, read.big.matrix . . . . .            | 20        |
| <b>Index</b>   | <b>23</b> |

---

|                   |  |
|-------------------|--|
| bigmemory-package | <i>Manage massive matrices with shared memory and memory-mapped files.</i> |
|-------------------|--|

---

Description

Create, store, access, and manipulate massive matrices. Matrices are, by default, allocated to shared memory and may use memory-mapped files. Packages **biganalytics**, **synchronicity**, **bigalgebra**, and **bigtabulate** provide advanced functionality. Access to and manipulation of a `big.matrix` object is exposed in R by an S4 class whose interface is similar to that of an R `matrix`. Use of these packages in parallel environments can provide substantial speed and memory efficiencies. **bigmemory** also provides a C++ framework for the development of new tools that can work both with `big.matrix` and native R `matrix` objects.

Details

|            |   |
|------------|---|
| Package:   | bigmemory   |
| Type:      | Package   |
| Version:   | 4.4.6   |
| Date:      | 2013-11-18  |
| License:   | LGPL-3  |
| Copyright: | (C) 2013 Michael J. Kane and John W. Emerson                    |
| URL:       | <a href="http://www.bigmemory.org">http://www.bigmemory.org</a> |
| LazyLoad:  | yes   |

Index of functions/methods (grouped in a friendly way):

`big.matrix`, `filebacked.big.matrix`, `as.big.matrix`

```

is.big.matrix, is.separated, is.filebacked

describe, attach.big.matrix, attach.resource

sub.big.matrix, is.sub.big.matrix

dim, dimnames, nrow, ncol, print, head, tail, typeof, length

read.big.matrix, write.big.matrix

mwhich

morder, mpermute

deepcopy

flush

```

Multi-gigabyte data sets challenge and frustrate R users, even on well-equipped hardware. Use of C/C++ can provide efficiencies, but is cumbersome for interactive data analysis and lacks the flexibility and power of R's rich statistical programming environment. The package **bigmemory** and sister packages **biganalytics**, **synchronicity**, **bigtabulate**, and **bigalgebra** bridge this gap, implementing massive matrices and supporting their manipulation and exploration. The data structures may be allocated to shared memory, allowing separate processes on the same computer to share access to a single copy of the data set. The data structures may also be file-backed, allowing users to easily manage and analyze data sets larger than available RAM and share them across nodes of a cluster. These features of the Bigmemory Project open the door for powerful and memory-efficient parallel analyses and data mining of massive data sets.

This project (**bigmemory** and its sister packages) is still actively developed, although the design and current features can be viewed as "stable." Please feel free to email us with any questions: [bigmemoryauthors@gmail.com](mailto:bigmemoryauthors@gmail.com).

## Note

Various options are available. `options(bigmemory.typecast.warning)` can be set to avoid annoying warnings that might occur if, for example, you assign R objects (typically type double) to char, short, or integer `big.matrix` objects. `options(bigmemory.print.warning)` protects against extracting and printing a massive matrix (which would involve the creation of a second massive copy of the matrix). `options(bigmemory.allow.dimnames)` by default prevents the setting of `dimnames` attributes, because they aren't allocated to shared memory and changes will not be visible across processes. `options(bigmemory.default.type)` is "double" by default (a change in default behavior as of 4.1.1) but may be changed by the user.

Versions  $\geq 4.0$  represent a major redesign, with the mutexes (locking) abstracted to package **synchronicity**, the exploratory data analysis functionality relocated to package **biganalytics**, and new linear algebra support available in package **bigalgebra**. Package **bigtabulate** extends the **bigmemory** package with table-, tapply-, and split-like behavior. The functions may also be used with regular R matrices for speed and memory-efficiency gains. Package **bigmemory** itself is now minimalist, providing only the core functionality. As an example, the `apply()` method appears in

**biganalytics**, supporting exploration and analysis, while **mwhich**, **morder** and **mpermute** appear in **bigmemory** as fundamental tools for data manipulation.

Versions <4.0 supported a limited number of columns (due to mutex limitations): roughly 50,000 on a typical Linux system. This restriction has been removed in versions >=4.0. There were row limitations (due to a bug that has now been fixed) in versions <=3.8 of roughly 1 billion, but this has been fixed in versions >=3.82.

Note that you can't simply use a `big.matrix` with many (most) existing R functions (e.g. `lm`, `kmeans`). One nice exception is `split`, because this function only accesses subsets of the matrix.

### Author(s)

Michael J. Kane and John W. Emerson

Maintainers: Michael J. Kane <bigmemoryauthors@gmail.com>

### References

The Bigmemory Project: <http://www.bigmemory.org/>.

### See Also

For example, `big.matrix`, `mwhich`, `read.big.matrix`

### Examples

```
# Our examples are all trivial in size, rather than burning huge amounts
# of memory.

x <- big.matrix(5, 2, type="integer", init=0,
               dimnames=list(NULL, c("alpha", "beta")))
x
x[1:2,]
x[,1] <- 1:5
x[, "alpha"]
colnames(x)
options(bigmemory.allow.dimnames=TRUE)
colnames(x) <- NULL
x[,]
```

---

as.big.matrix-methods *Create a "big.matrix" from a matrix or vector.*

---

### Description

Create a `big.matrix` from a matrix or vector or `data.frame`; a vector will result in a `big.matrix` with one column. A data frame will have character vectors converted to factors, and then all factors converted to numeric factor levels. All labels or character values will be lost.

**Methods**

```
signature(x = "matrix") ...
signature(x = "vector") ...
signature(x = "data.frame") ...
```

---

```
big.matrix, filebacked.big.matrix, etc...
```

*The core “big.matrix” operations.*

---

**Description**

Create a big.matrix (or check to see if an object is a big.matrix, or create a big.matrix from a [matrix](#), and so on). The big.matrix may be file-backed.

**Usage**

```
big.matrix(nrow, ncol, type = options()$bigmemory.default.type,
  init = NULL, dimnames = NULL, separated = FALSE,
  backingfile = NULL, backingpath = NULL, descriptorfile = NULL,
  binarydescriptor=FALSE, shared = TRUE)
filebacked.big.matrix(nrow, ncol,
  type = options()$bigmemory.default.type, init = NULL,
  dimnames = NULL, separated = FALSE, backingfile = NULL,
  backingpath = NULL, descriptorfile = NULL, binarydescriptor=FALSE)
as.big.matrix(x, type = NULL, separated = FALSE,
  backingfile = NULL, backingpath = NULL, descriptorfile = NULL,
  binarydescriptor=FALSE, shared=TRUE)
is.big.matrix(x)
is.separated(x)
is.filebacked(x)
is.shared(x)
is.readonly(x)
is.nil(address)
```

**Arguments**

|      |  |
|------|--|
| x    | a matrix, vector, or data.frame for as.big.matrix; if a vector, a one-column big.matrix is created by as.big.matrix; if a data.frame, see details. For the is.* functions, x is likely a big.matrix. |
| nrow | number of rows.  |
| ncol | number of columns.   |
| type | the type of the atomic element (options()\$bigmemory.default.type by default – “double” – but can be changed by the user to “integer”, “short”, or “char”).  |

|                               |   |
|-------------------------------|---|
| <code>init</code>             | a scalar value for initializing the matrix (NULL by default to avoid unnecessary time spent doing the initializing).  |
| <code>dimnames</code>         | a list of the row and column names; use with caution for large objects.   |
| <code>separated</code>        | use separated column organization of the data; see details.   |
| <code>backingfile</code>      | the root name for the file(s) for the cache of <code>x</code> .   |
| <code>backingpath</code>      | the path to the directory containing the file backing cache.  |
| <code>descriptorfile</code>   | the name of the file to hold the backingfile description, for subsequent use with <a href="#">attach.big.matrix</a> ; if NULL, the backingfile is used as the root part of the descriptor file name. The descriptor file is placed in the same directory as the backing files.  |
| <code>binarydescriptor</code> | the flag to specify if the binary RDS format should be used for the backingfile description, for subsequent use with <a href="#">attach.big.matrix</a> ; if NULL or FALSE, the <code>dput()</code> file format is used.   |
| <code>shared</code>           | TRUE by default, and always TRUE if the <code>big.matrix</code> is file-backed. For a non-filebacked <code>big.matrix</code> , <code>shared=FALSE</code> uses non-shared memory, which can be more stable for large (say, >50% of RAM) objects. Shared memory allocation can sometimes fail in such cases due to exhausted shared-memory resources in the system. |
| <code>address</code>          | an <code>externalptr</code> , so <code>is.nil(x@address)</code> might be a sensible thing to want to check, but it's pretty obscure.  |

## Details

A `big.matrix` consists of an object in R that does nothing more than point to the data structure implemented in C++. The object acts much like a traditional R matrix, but helps protect the user from many inadvertant memory-consuming pitfalls of traditional R matrices and data frames.

There are two `big.matrix` types which manage data in different ways. A standard, shared `big.matrix` is constrained to available RAM, and may be shared across separate R processes. A file-backed `big.matrix` may exceed available RAM by using hard drive space, and may also be shared across processes. The atomic types of these matrices may be double, integer, short, or char (8, 4, 2, and 1 bytes, respectively).

If `x` is a `big.matrix`, then `x[1:5,]` is returned as an R matrix containing the first five rows of `x`. If `x` is of type double, then the result will be numeric; otherwise, the result will be an integer R matrix. The expression `x` alone will display information about the R object (e.g. the external pointer) rather than evaluating the matrix itself (the user should try `x[,]` with extreme caution, recognizing that a huge R matrix will be created).

If `x` has a huge number of rows and/or columns, then the use of `rownames` and/or `colnames` will be extremely memory-intensive and should be avoided. If `x` has a huge number of columns and `separated=TRUE` is used (this isn't typically recommended), the user might want to store the transpose as there is overhead of a pointer for each column in the matrix. If `separated` is TRUE, then the memory is allocated into separate vectors for each column. Use this option with caution if you have a large number of columns, as shared-memory segments are limited by OS and hardware combinations. If `separated` is FALSE, the matrix is stored in traditional column-major format. The function `is.separated()` returns the separation type of the `big.matrix`.

When a `big.matrix`, `x`, is passed as an argument to a function, it is essentially providing call-by-reference rather than call-by-value behavior. If the function modifies any of the values of `x`, the changes are not limited in scope to a local copy within the function. This introduces the possibility of side-effects, in contrast to standard R behavior.

A file-backed `big.matrix` may exceed available RAM in size by using a file cache (or possibly multiple file caches, if `separated=TRUE`). This can incur a substantial performance penalty for such large matrices, but less of a penalty than most other approaches for handling such large objects. A side-effect of creating a file-backed object is not only the file-backing(s), but a descriptor file (in the same directory) that is needed for subsequent attachments (see [attach.big.matrix](#)).

Note that we do not allow setting or changing the `dimnames` attributes by default; such changes would not be reflected in the descriptor objects or in shared memory. To override this, set `options(bigmemory.allow.dimnames=TRUE)`.

It should also be noted that a user can create an “anonymous” file-backed `big.matrix` by specifying `""` as the `filebacking` argument. In this case, the backing resides in the temporary directory and a descriptor file is not created. These should be used with caution since even anonymous backings use disk space which could eventually fill the hard drive. Anonymous backings are removed either manually, by a user, or automatically, when the operating system deems it appropriate.

Finally, note that `as.big.matrix` can coerce data frames. It does this by making any character columns into factors, and then making all factors numeric before forming the `big.matrix`. Level labels are not preserved and must be managed by the user if desired.

## Value

A `big.matrix` is returned (for `big.matrix` and `filebacked.big.matrix`, and `as.big.matrix`), and `TRUE` or `FALSE` for `is.big.matrix` and the other functions.

## Author(s)

John W. Emerson and Michael J. Kane <bigmemoryauthors@gmail.com>

## References

The Bigmemory Project: <http://www.bigmemory.org/>.

## See Also

[bigmemory](#), and perhaps the class documentation of [big.matrix](#); [attach.big.matrix](#) and [describe](#). Sister packages **biganalytics**, **bigtabulate**, **synchronicity**, and **bigalgebra** provide advanced functionality.

## Examples

```
x <- big.matrix(10, 2, type='integer', init=-5)
options(bigmemory.allow.dimnames=TRUE)
colnames(x) <- c("alpha", "beta")
is.big.matrix(x)
dim(x)
colnames(x)
```

```

rownames(x)
x[,]
x[1:8,1] <- 11:18
colnames(x) <- NULL
x[,]

x <- as.big.matrix(matrix(-5, 10, 2))
colnames(x) <- c("alpha", "beta")
is.big.matrix(x)
dim(x)
colnames(x)
rownames(x)
x[1:8,1] <- 11:18
x[,]

# The following shared memory example is quite silly, as you wouldn't
# likely do this in a single R session. But if zdescription were
# passed to another R session via SNOW, foreach, or even by a
# simple file read/write, then the attach.big.matrix() within the
# second R process would give access to the same object in memory.
# Please see the package vignette for real examples.

z <- big.matrix(3, 3, type='integer', init=3)
z[,]
dim(z)
z[1,1] <- 2
z[,]
zdescription <- describe(z)
zdescription
y <- attach.big.matrix(zdescription)
y[,]
y
z
y[1,1] <- -100
y[,]
z[,]

# A short filebacked example, showing the creation of associated files:
files <- dir()
files[grep("example.bin", files)]
z <- filebacked.big.matrix(3, 3, type='integer', init=123,
                           backingfile="example.bin",
                           descriptorfile="example.desc",
                           dimnames=list(c('a','b','c'), c('d', 'e', 'f')))

z[,]
files <- dir()
files[grep("example.bin", files)]
zz <- attach.big.matrix("example.desc")
zz[,]
zz[1,1] <- 0
zzz <- attach.big.matrix(describe(z))
zzz[,]
```



```
is.nil(z@address)
```

---

|                  |                    |
|------------------|--------------------|
| big.matrix-class | Class "big.matrix" |
|------------------|--------------------|

---

## Description

The `big.matrix` class is designed for matrices with elements of type double, integer, short, or char. A `big.matrix` acts much like a traditional R matrix, but helps protect the user from many inadvertant memory-consuming pitfalls of traditional R matrices and data frames. The objects are allocated to shared memory, and if file-backing is used they may exceed virtual memory in size. Sadly, 32-bit operating system constraints – largely Windows and some MacOS versions – will be a limiting factor with file-backed matrices; 64-bit operating systems are recommended.

## Objects from the Class

Unlike many R objects, objects should not be created by calls of the form `new("big.matrix", ...)`. The functions `big.matrix()` and `filebacked.big.matrix()` are intended for the user.

## Slots

**address:** Object of class "externalptr" points to the memory location of the C++ data structure.

## Methods

As you would expect:

```
signature(x = "big.matrix", i = "ANY", j = "ANY"): ...
[<-signature(x = "big.matrix", i = "ANY", j = "missing"): ...
[<- signature(x = "big.matrix", i = "missing", j = "ANY"): ...
[<- signature(x = "big.matrix", i = "missing", j = "missing"): ...
[<- signature(x = "big.matrix", i = "matrix", j = "missing"): ...
[ signature(x = "big.matrix", i = "ANY", j = "ANY", drop = "missing"): ...
[ signature(x = "big.matrix", i = "ANY", j = "ANY", drop = "logical"): ...
[ signature(x = "big.matrix", i = "ANY", j = "missing", drop = "missing"): ...
[ signature(x = "big.matrix", i = "ANY", j = "missing", drop = "logical"): ...
[ signature(x = "big.matrix", i = "matrix", j = "missing", drop = "logical"): ...
[ signature(x = "big.matrix", i = "missing", j = "ANY", drop = "missing"): ...
[ signature(x = "big.matrix", i = "missing", j = "ANY", drop = "logical"): ...
[ signature(x = "big.matrix", i = "missing", j = "missing", drop = "missing"): ...
[ signature(x = "big.matrix", i = "missing", j = "missing", drop = "logical"): ...
```

The following are probably more interesting:

**describe** signature(x = "big.matrix"): provide necessary and sufficient information for the sharing or re-attaching of the object.

**dim** signature(x = "big.matrix"): returns the dimension of the big.matrix.

**length** signature(x = "big.matrix"): returns the product of the dimensions of the big.matrix.

**dimnames<-** signature(x = "big.matrix", value = "list"): set the row and column names, prohibited by default (see [bigmemory](#) to override).

**dimnames** signature(x = "big.matrix"): get the row and column names.

**head** signature(x = "big.matrix"): get the first 6 (or n) rows.

**as.matrix** signature(x = "big.matrix"): coerce a big.matrix to a matrix.

**is.big.matrix** signature(x = "big.matrix"): return TRUE if it's a big.matrix.

**is.filebacked** signature(x = "big.matrix"): return TRUE if there is a file-backing.

**is.separated** signature(x = "big.matrix") : return TRUE if the big.matrix is organized as a separated column vectors.

**is.sub.big.matrix** signature(x = "big.matrix"): return TRUE if this is a sub-matrix of a big.matrix.

**ncol** signature(x = "big.matrix"): returns the number of columns.

**nrow** signature(x = "big.matrix"): returns the number of rows.

**print** signature(x = "big.matrix"): a traditional print() is intentionally disabled, and returns head(x) unless options()\$bm.print.warning==FALSE; in this case, print(x[,]) is the result, which could be very big!

**sub.big.matrix** signature(x = "big.matrix"): for contiguous submatrices.

**tail** signature(x = "big.matrix"): returns the last 6 (or n) rows.

**typeof** signature(x = "big.matrix"): return the type of the atomic elements of the big.matrix.

**write.big.matrix** signature(bigMat = "big.matrix", fileName = "character"): produce an ASCII file from the big.matrix.

**apply** signature(x = "big.matrix"): apply() where MARGIN may only be 1 or 2, but otherwise conforming to what you would expect from apply().

## See Also

[big.matrix](#)

## Examples

```
showClass("big.matrix")
```

---

```
big.matrix.descriptor-class  
      Class "big.matrix.descriptor"
```

---

## Description

An object of this class contains necessary and sufficient information to “attach” a shared or file-backed [big.matrix](#).

## Objects from the Class

Objects should not be created by calls of the form `new("big.matrix.descriptor", ...)`, but should use the [describe](#) function.

## Slots

**description:** Object of class “list”; details omitted.

## Extends

Class “[descriptor](#)”, directly.

## Methods

```
attach.resource signature(obj = "big.matrix.descriptor"): ...  
sub.big.matrix signature(x = "big.matrix.descriptor"): ...
```

## Note

We provide `attach.resource` for convenience, but expect most users will prefer [attach.big.matrix](#).

## References

Other types of descriptors are defined in package **synchronicity**.

## See Also

See also [attach.big.matrix](#).

## Examples

```
showClass("big.matrix.descriptor")
```

---

deepcopy

*Produces a physical copy of a “big.matrix”*


---

### Description

This is needed to make a duplicate of a `big.matrix`, with the new copy optionally filebacked.

### Usage

```
deepcopy(x, cols = NULL, rows = NULL, y = NULL, type = NULL,
         separated = NULL, backingfile = NULL, backingpath = NULL,
         descriptorfile = NULL, binarydescriptor=FALSE, shared=TRUE)
```

### Arguments

|                               |   |
|-------------------------------|---|
| <code>x</code>                | a <code>big.matrix</code> .   |
| <code>cols</code>             | possible subset of columns for the deepcopy; could be numeric, named, or logical.   |
| <code>rows</code>             | possible subset of rows for the deepcopy; could be numeric, named, or logical.  |
| <code>y</code>                | optional destination object (matrix or <code>big.matrix</code> ); if not specified, a <code>big.matrix</code> will be created.  |
| <code>type</code>             | preferably specified, “integer” for example.  |
| <code>separated</code>        | use separated column organization of the data instead of column-major organization; use with caution if the number of columns is large.   |
| <code>backingfile</code>      | the root name for the file(s) for the cache of <code>x</code> .   |
| <code>backingpath</code>      | the path to the directory containing the file-backing cache.  |
| <code>descriptorfile</code>   | we recommend specifying this for file-backing.  |
| <code>binarydescriptor</code> | the flag to specify if the binary RDS format should be used for the backingfile description, for subsequent use with <code>attach.big.matrix</code> ; if NULL or FALSE, the <code>dput()</code> file format is used.  |
| <code>shared</code>           | TRUE by default, and always TRUE if the <code>big.matrix</code> is file-backed. For a non-filebacked <code>big.matrix</code> , <code>shared=FALSE</code> uses non-shared memory, which can be more stable for large (say, >50% of RAM) objects. Shared memory allocation can sometimes fail in such cases due to exhausted shared-memory resources in the system. |

### Details

This is needed to make a duplicate of a `big.matrix`, because traditional R syntax would only copy the R object (the pointer to the `big.matrix` rather than the `big.matrix` itself). It can also make a copy of only a subset of columns.

**Value**

a `big.matrix`.

**See Also**

`big.matrix`

**Examples**

```
x <- as.big.matrix(matrix(1:30, 10, 3))
y <- deepcopy(x, -1)    # Don't include the first column.
x
y
head(x)
head(y)
```

---

describe, attach.big.matrix, attach.resource

*The basic “big.matrix” operations for sharing and re-attaching.*

---

**Description**

The describe function returns the information needed by attach.big.matrix to reference a shared or file-backed big.matrix object. The attach.big.matrix and attach.resource functions create a new big.matrix object based on the descriptor information referencing previously allocated shared-memory or file-backed matrices.

**Usage**

```
describe(x)
attach.big.matrix(obj, ...)
attach.resource(obj, ...)
```

**Arguments**

|                  |   |
|------------------|---|
| <code>x</code>   | a <code>big.matrix</code> object  |
| <code>obj</code> | an object as returned by describe() or, optionally, the filename of the descriptor for a filebacked matrix, assumed to be in the directory specified by the path (if one is provided) |
| <code>...</code> | possibly path which gives the path where the descriptor and/or filebacking can be found   |

**Details**

The describe function returns a list of the information needed to attach to a big.matrix object. A descriptor file is automatically created when a new filebacked big.matrix is created.

**Value**

describe returns a list of the information needed to attach to a `big.matrix` object.

`attach.big.matrix` return a new instance of type `big.matrix` corresponding to a shared-memory or file-backed `big.matrix`.

**Author(s)**

Michael J. Kane and John W. Emerson <bigmemoryauthors@gmail.com>

**See Also**

[bigmemory](#), [big.matrix](#), or the class documentation [big.matrix](#).

**Examples**

```
# The example is quite silly, as you wouldn't likely do this in a
# single R session. But if zdescription were passed to another R session
# via SNOW, foreach, or even by a simple file read/write,
# then the attach of the second R process would give access to the
# same object in memory. Please see the package vignette for real examples.
z <- big.matrix(3, 3, type='integer', init=3)
z[,]
dim(z)
z[1,1] <- 2
z[,]
zdescription <- describe(z)
zdescription
y <- attach.big.matrix(zdescription)
y[,]
y
z
zz <- attach.resource(zdescription)
zz[1,1] <- -100
y[,]
z[,]
```

---

flush

---

*Updating a big.matrix filebacking.*


---

**Description**

For a file-backed `big.matrix` object, `flush()` forces any modified information to be written to the file-backing.

**Usage**

```
flush(con)
```

**Arguments**

con                    a filebacked `big.matrix`.

**Details**

This function flushes any modified data (in RAM) of a file-backed `big.matrix` to disk. This may be useful for improving performance in cases where allowing the operating system to decide on flushing creates a bottleneck (likely near the threshold of available RAM).

**Value**

TRUE or FALSE (invisible), indicating whether or not the flush was successful.

**Examples**

```
x <- big.matrix(nrow=3, ncol=3, backingfile='flushtest.bin',
  descriptorfile='flushtest.desc', type='integer')
x[1,1] <- 0
flush(x)
```

---

|                  |   |
|------------------|---|
| morder, mpermute | <i>Ordering and row-permuting functions for “big.matrix” and “matrix” objects</i> |
|------------------|---|

---

**Description**

The `morder` function returns a permutation of row indices which can be used to rearrange an object according to the values in the specified columns (a multi-column ordering). The `mpermute` function actually reorders the rows of a `big.matrix` or `matrix` based on an order vector or a desired ordering on a set of columns.

**Usage**

```
morder(x, cols, na.last=TRUE, decreasing = FALSE)
mpermute(x, order=NULL, cols=NULL, allow.duplicates=FALSE, ...)
```

**Arguments**

|            |  |
|------------|--|
| x          | A <code>big.matrix</code> or <code>matrix</code> object with numeric values.   |
| cols       | The columns of x to get the ordering for or reorder on   |
| na.last    | for controlling the treatment of NAs. If TRUE, missing values in the data are put last; if FALSE, they are put first; if NA, they are removed. |
| decreasing | logical. Should the sort order be increasing or decreasing?  |
| order      | A vector specifying the reordering of rows, i.e. the result of a call to <code>order</code> or <code>morder</code> .                           |

`allow.duplicates`      ff TRUE, allows a row to be duplicated in the resulting `big.matrix` or `matrix` (i.e. in this case, `order` would not need to be a permutation of `1:nrow(x)`).

`...`                    optional parameters to pass to `morder` when `cols` is specified instead of just using `order`.

## Details

The `morder` function behaves similar to `order`, returning a permutation of `1:nrow(x)` which rearranges objects according to the values in the specified columns. However, `morder` takes a `big.matrix` or an `R matrix` (with numeric type) and a set of columns (`cols`) with which to determine the ordering; `morder` does not incur the same memory overhead required by `order`, and runs more quickly.

The `mpermute` function changes the row ordering of a `big.matrix` or `matrix` based on a vector order or an ordering based on a set of columns specified by `cols`. It should be noted that this function has side-effects, that is `x` is changed when this function is called.

## Value

`morder` returns an ordering vector.

`mpermute` returns nothing but does change the contents of `x`. This type of a side-effect is generally frowned upon in R, but we “break” the rules here to avoid memory overhead and improve performance.

## Author(s)

Michael J. Kane <bigmemoryauthors@gmail.com>

## See Also

[order](#)

## Examples

```
m = matrix(as.double(as.matrix(iris)), nrow=nrow(iris))
morder(m, 1)
order(m[,1])

m[order(m[,1]), 2]
mpermute(m, cols=1)
m[,2]
```



---

|        |   |
|--------|---|
| mwhich | <i>Expanded “which”-like functionality.</i> |
|--------|---|

---

## Description

Implements [which](#)-like functionality for a [big.matrix](#), with additional options for efficient comparisons (executed in C++); also works for regular numeric matrices without the memory overhead.

## Usage

```
mwhich(x, cols, vals, comps, op = 'AND')
```

## Arguments

|       |  |
|-------|--|
| x     | a <a href="#">big.matrix</a> (or a numeric matrix; see below).   |
| cols  | a vector of column indices or names.   |
| vals  | a list (one component for each of cols) of vectors of length 1 or 2; length 1 is used to test equality (or inequality), while vectors of length 2 are used for checking values in the range ( $-\text{Inf}$ and $\text{Inf}$ are allowed). If a scalar or vector of length 2 is provided instead of a list, it will be replicated <code>length(cols)</code> times. |
| comps | a list of operators (one component for each of cols), including 'eq', 'neq', 'le', 'lt', 'ge' and 'gt'. If a single operator, it will be replicated <code>length(cols)</code> times.   |
| op    | the comparison operator for combining the results of the individual tests, either 'AND' or 'OR'.   |

## Details

To improve performance and avoid the creation of massive temporary vectors in R when doing comparisons, `mwhich()` efficiently executes column-by-column comparisons of values to the specified values or ranges, and then returns the row indices satisfying the comparison specified by the `op` operator. More advanced comparisons are then possible (and memory-efficient) in R by doing set operations ([union](#) and [intersect](#), for example) on the results of multiple `mwhich()` calls.

Note that NA is a valid argument in conjunction with 'eq' or 'neq', replacing traditional `is.na()` calls. And both  $-\text{Inf}$  and  $\text{Inf}$  can be used for one-sided inequalities.

If `mwhich()` is used with a regular numeric R matrix, we access the data directly and thus incur no memory overhead. Interested developers might want to look at our code for this case, which uses a handy pointer trick (accessor) in C++.

## Value

a vector of row indices satisfying the criteria.

## Author(s)

John W. Emerson <[bigmemoryauthors@gmail.com](mailto:bigmemoryauthors@gmail.com)>

**See Also**

[big.matrix](#), [which](#)

**Examples**

```
x <- as.big.matrix(matrix(1:30, 10, 3))
options(bigmemory.allow.dimnames=TRUE)
colnames(x) <- c("A", "B", "C")
x[,]
x[mwhich(x, 1:2, list(c(2,3), c(11,17)),
               list(c('ge','le'), c('gt', 'lt'))), 'OR'),]

x[mwhich(x, c("A","B"), list(c(2,3), c(11,17)),
               list(c('ge','le'), c('gt', 'lt'))), 'AND'),]

# These should produce the same answer with a regular matrix:
y <- matrix(1:30, 10, 3)
y[mwhich(y, 1:2, list(c(2,3), c(11,17)),
          list(c('ge','le'), c('gt', 'lt'))), 'OR'),]

y[mwhich(y, -3, list(c(2,3), c(11,17)),
          list(c('ge','le'), c('gt', 'lt'))), 'AND'),]

x[1,1] <- NA
mwhich(x, 1:2, NA, 'eq', 'OR')
mwhich(x, 1:2, NA, 'neq', 'AND')

# Column 1 equal to 4 and/or column 2 less than or equal to 16:
mwhich(x, 1:2, list(4, 16), list('eq', 'le'), 'OR')
mwhich(x, 1:2, list(4, 16), list('eq', 'le'), 'AND')

# Column 2 less than or equal to 15:
mwhich(x, 2, 15, 'le')

# No NAs in either column, and column 2 strictly less than 15:
mwhich(x, c(1:2,2), list(NA, NA, 15), list('neq', 'neq', 'lt'), 'AND')

x <- big.matrix(4, 2, init=1, type="double")
x[1,1] <- Inf
mwhich(x, 1, Inf, 'eq')
mwhich(x, 1, 1, 'gt')
mwhich(x, 1, 1, 'le')
```

**Description**

Implements [which](#)-like functionality for a [big.matrix](#), with additional options for efficient comparisons (executed in C++); also works for regular numeric matrices without the memory overhead.

**Methods**

```
signature(x = "big.matrix", cols = "ANY", vals = "ANY", comps = "ANY", op = "character")
...
signature(x = "big.matrix", cols = "ANY", vals = "ANY", comps = "ANY", op = "missing")
...
signature(x = "matrix", cols = "ANY", vals = "ANY", comps = "ANY", op = "character")
...
signature(x = "matrix", cols = "ANY", vals = "ANY", comps = "ANY", op = "missing")
...
```

**See Also**

[big.matrix](#), [which](#), [mwhich](#)

---

sub.big.matrix, is.sub.big.matrix  
*Submatrix support.*

---

**Description**

This doesn't create a copy, it just provides a new version of the class which provides behavior for a contiguous submatrix of the [big.matrix](#). Non-contiguous submatrices are not supported.

**Usage**

```
sub.big.matrix(x, firstRow = 1, lastRow = NULL,
               firstCol = 1, lastCol = NULL, backingpath='')
is.sub.big.matrix(x)
```

**Arguments**

|             |  |
|-------------|--|
| x           | either a <a href="#">big.matrix</a> or a descriptor.   |
| firstRow    | the first row of the submatrix.                        |
| lastRow     | the last row of the submatrix if not NULL.             |
| firstCol    | the first column of the submatrix.                     |
| lastCol     | the last column of the submatrix if not NULL.          |
| backingpath | required path to the filebacked object, if applicable. |

## Details

The `sub.big.matrix` function allows a user to create a `big.matrix` object that references a contiguous set of columns and rows of another `big.matrix` object.

The `is.sub.big.matrix` function returns `TRUE` if the specified argument is a `sub.big.matrix` object and return `FALSE` otherwise.

## Value

A `big.matrix` which is actually a submatrix of a larger `big.matrix`. It is not a physical copy. Only contiguous blocks may form a submatrix.

## See Also

`big.matrix`

## Examples

```
x <- big.matrix(10, 5, init=0, type="double")
x[,] <- 1:50
y <- sub.big.matrix(x, 2, 9, 2, 3)
y[,]
y[1,1] <- -99
x[,]
```

---

`write.big.matrix, read.big.matrix`

*File interface for a “big.matrix”*

---

## Description

Create a `big.matrix` by reading from a suitably-formatted ASCII file, or write the contents of a `big.matrix` to a file.

## Usage

```
write.big.matrix(x, filename, row.names = FALSE,
                 col.names = FALSE, sep=',')
read.big.matrix(filename, sep = ',', header = FALSE,
                 col.names = NULL, row.names = NULL,
                 has.row.names=FALSE, ignore.row.names=FALSE,
                 type = NA, skip = 0, separated = FALSE,
                 backingfile = NULL, backingpath = NULL,
                 descriptorfile = NULL, binarydescriptor=FALSE,
                 extraCols = NULL, shared=TRUE)
```

## Arguments

|                  |  |
|------------------|--|
| x                | a <a href="#">big.matrix</a> .   |
| filename         | the name of an input/output file.  |
| sep              | a field delimiter.   |
| header           | if TRUE, the first line (after a possible skip) should contain column names.   |
| col.names        | a vector of names, use them even if column names exist in the file.  |
| row.names        | a vector of names, use them even if row names appear to exist in the file.   |
| has.row.names    | if TRUE, then the first column contains row names.   |
| ignore.row.names | if TRUE when has.row.names==TRUE, the row names will be ignored.   |
| type             | preferably specified, "integer" for example.   |
| skip             | number of lines to skip at the head of the file.   |
| separated        | use separated column organization of the data instead of column-major organization.  |
| backingfile      | the root name for the file(s) for the cache of x.  |
| backingpath      | the path to the directory containing the file backing cache.   |
| descriptorfile   | the file to be used for the description of the filebacked matrix.  |
| binarydescriptor | the flag to specify if the binary RDS format should be used for the backingfile description, for subsequent use with <a href="#">attach.big.matrix</a> ; if NULL or FALSE, the dput() file format is used. |
| extraCols        | the optional number of extra columns to be appended to the matrix for future use.  |
| shared           | if TRUE, the resulting big.matrix can be shared across processes.  |

## Details

Files must contain only one atomic type (all integer, for example). You, the user, should know whether your file has row and/or column names, and various combinations of options should be helpful in obtaining the desired behavior.

When reading from a file, if type is not specified we try to make a reasonable guess for you without making any guarantees at this point. Unless you have really large integer values, we recommend you consider "short". If you have something that is essentially categorical, you might even be able use "char", with huge memory savings for large data sets.

Any non-numeric entry will be ignored and replaced with NA, so reading something that traditionally would be a data.frame won't cause an error. A warning is issued.

Wishlist: we'd like to provide an option to ignore specified columns while doing reads. Or perhaps to specify columns targeted for factor or character conversion to numeric values. Would you use such features? Email us and let us know!

## Value

a [big.matrix](#) object is returned by read.big.matrix, while write.big.matrix creates an output file (a path could be part of filename).

**Author(s)**

John W. Emerson and Michael J. Kane <bigmemoryauthors@gmail.com>

**See Also**

[big.matrix](#)

**Examples**

```
# Without specifying the type, this big.matrix x will hold integers.
x <- as.big.matrix(matrix(1:10, 5, 2))
x[2,2] <- NA
x[, ]
write.big.matrix(x, "foo.txt")

# Just for fun, I'll read it back in as character (1-byte integers):
y <- read.big.matrix("foo.txt", type="char")
y[, ]

# Other examples:
w <- as.big.matrix(matrix(1:10, 5, 2), type='double')
w[1,2] <- NA
w[2,2] <- -Inf
w[3,2] <- Inf
w[4,2] <- NaN
w[, ]
write.big.matrix(w, "bar.txt")
w <- read.big.matrix("bar.txt", type="double")
w[, ]
w <- read.big.matrix("bar.txt", type="short")
w[, ]

# Another example using row names (which we don't like).
x <- as.big.matrix(as.matrix(iris), type='double')
rownames(x) <- as.character(1:nrow(x))
head(x)
write.big.matrix(x, 'IrisData.txt', col.names=TRUE, row.names=TRUE)
y <- read.big.matrix("IrisData.txt", header=TRUE, has.row.names=TRUE)
head(y)

# The following would fail with a dimension mismatch:
if (FALSE) y <- read.big.matrix("IrisData.txt", header=TRUE)
```

# Index

## \*Topic **classes**

- big.matrix,
  - filebacked.big.matrix, etc...,  
5
- big.matrix-class, 9
- big.matrix.descriptor-class, 11
- describe, attach.big.matrix,
  - attach.resource, 13

## \*Topic **methods**

- as.big.matrix-methods, 4
- big.matrix,
  - filebacked.big.matrix, etc...,  
5
- deepcopy, 12
- describe, attach.big.matrix,
  - attach.resource, 13
- flush, 14
- mwhich, 17
- mwhich-methods, 18
- sub.big.matrix, is.sub.big.matrix,  
19
- write.big.matrix, read.big.matrix,  
20

## \*Topic **package**

- bigmemory-package, 2
- [,big.matrix,ANY,ANY,logical-method  
(big.matrix-class), 9
- [,big.matrix,ANY,ANY,missing-method  
(big.matrix-class), 9
- [,big.matrix,ANY,missing,logical-method  
(big.matrix-class), 9
- [,big.matrix,ANY,missing,missing-method  
(big.matrix-class), 9
- [,big.matrix,matrix,missing,missing-method  
(big.matrix-class), 9
- [,big.matrix,missing,ANY,logical-method  
(big.matrix-class), 9
- [,big.matrix,missing,ANY,missing-method  
(big.matrix-class), 9

- [,big.matrix,missing,missing,logical-method  
(big.matrix-class), 9
- [,big.matrix,missing,missing,missing-method  
(big.matrix-class), 9
- [.big.matrix(big.matrix,
  - filebacked.big.matrix,  
etc...), 5
- [<-,big.matrix,ANY,ANY-method  
(big.matrix-class), 9
- [<-,big.matrix,ANY,missing-method  
(big.matrix-class), 9
- [<-,big.matrix,matrix,missing-method  
(big.matrix-class), 9
- [<-,big.matrix,missing,ANY-method  
(big.matrix-class), 9
- [<-,big.matrix,missing,missing-method  
(big.matrix-class), 9

- as.big.matrix(big.matrix,
  - filebacked.big.matrix,  
etc...), 5
- as.big.matrix,data.frame-method  
(as.big.matrix-methods), 4
- as.big.matrix,matrix-method  
(as.big.matrix-methods), 4
- as.big.matrix,vector-method  
(as.big.matrix-methods), 4
- as.big.matrix-methods, 4
- as.matrix, big.matrix-method  
(big.matrix-class), 9
- attach.big.matrix, 6, 7, 11, 12, 21
- attach.big.matrix(describe,
  - attach.big.matrix,  
attach.resource), 13
- attach.resource(describe,
  - attach.big.matrix,  
attach.resource), 13
- attach.resource, big.matrix.descriptor-method  
(big.matrix.descriptor-class),  
11

- attach.resource, character-method  
(big.matrix.descriptor-class),  
11
- big.matrix, 2-4, 7, 10-15, 17-22
- big.matrix (big.matrix,  
filebacked.big.matrix,  
etc...), 5
- big.matrix, filebacked.big.matrix,  
etc..., 5
- big.matrix-class, 9
- big.matrix.descriptor-class, 11
- bigmemory, 7, 10, 14
- bigmemory (bigmemory-package), 2
- bigmemory-package, 2
- data.frame, 4
- deepcopy, 12
- describe, 7, 11
- describe (describe, attach.big.matrix,  
attach.resource), 13
- describe, attach.big.matrix,  
attach.resource, 13
- describe, big.matrix-method  
(big.matrix-class), 9
- descriptor, 11
- descriptor-class  
(big.matrix.descriptor-class),  
11
- dim, big.matrix-method  
(big.matrix-class), 9
- dimnames, big.matrix-method  
(big.matrix-class), 9
- dimnames<-, big.matrix, list-method  
(big.matrix-class), 9
- filebacked.big.matrix (big.matrix,  
filebacked.big.matrix,  
etc...), 5
- flush, 14
- flush, big.matrix-method  
(big.matrix-class), 9
- head, big.matrix-method  
(big.matrix-class), 9
- intersect, 17
- is.big.matrix (big.matrix,  
filebacked.big.matrix,  
etc...), 5
- is.big.matrix, ANY-method  
(big.matrix-class), 9
- is.big.matrix, big.matrix-method  
(big.matrix-class), 9
- is.filebacked (big.matrix,  
filebacked.big.matrix,  
etc...), 5
- is.filebacked, big.matrix-method  
(big.matrix-class), 9
- is.nil (big.matrix,  
filebacked.big.matrix,  
etc...), 5
- is.readonly (big.matrix,  
filebacked.big.matrix,  
etc...), 5
- is.readonly, big.matrix-method  
(big.matrix-class), 9
- is.separated (big.matrix,  
filebacked.big.matrix,  
etc...), 5
- is.separated, big.matrix-method  
(big.matrix-class), 9
- is.shared (big.matrix,  
filebacked.big.matrix,  
etc...), 5
- is.shared, big.matrix-method  
(big.matrix-class), 9
- is.sub.big.matrix (sub.big.matrix,  
is.sub.big.matrix), 19
- is.sub.big.matrix, big.matrix-method  
(big.matrix-class), 9
- kmeans, 4
- length, big.matrix-method  
(big.matrix-class), 9
- lm, 4
- matrix, 2, 5
- morder, 4
- morder (morder, mpermute), 15
- morder, mpermute, 15
- mpermute, 4
- mpermute (morder, mpermute), 15
- mwhich, 4, 17, 19
- mwhich, big.matrix, ANY, ANY, ANY, character-method  
(mwhich-methods), 18
- mwhich, big.matrix, ANY, ANY, ANY, missing-method  
(mwhich-methods), 18



`mwhich`, `matrix`, `ANY`, `ANY`, `ANY`, `character-method`  
    (`mwhich-methods`), 18  
`mwhich`, `matrix`, `ANY`, `ANY`, `ANY`, `missing-method`  
    (`mwhich-methods`), 18  
`mwhich-methods`, 18  
  
`ncol`, `big.matrix-method`  
    (`big.matrix-class`), 9  
`nrow`, `big.matrix-method`  
    (`big.matrix-class`), 9  
  
`order`, 16  
  
`print`, `big.matrix-method`  
    (`big.matrix-class`), 9  
  
`read.big.matrix`, 4  
`read.big.matrix` (`write.big.matrix`,  
    `read.big.matrix`), 20  
`read.big.matrix`, `character-method`  
    (`big.matrix-class`), 9  
  
`split`, 4  
`sub.big.matrix` (`sub.big.matrix`,  
    `is.sub.big.matrix`), 19  
`sub.big.matrix`, `is.sub.big.matrix`, 19  
`sub.big.matrix`, `big.matrix-method`  
    (`big.matrix-class`), 9  
`sub.big.matrix`, `big.matrix.descriptor-method`  
    (`big.matrix.descriptor-class`),  
    11  
  
`tail`, `big.matrix-method`  
    (`big.matrix-class`), 9  
`typeof`, `big.matrix-method`  
    (`big.matrix-class`), 9  
  
`union`, 17  
  
`which`, 17–19  
`write.big.matrix` (`write.big.matrix`,  
    `read.big.matrix`), 20  
`write.big.matrix`, `read.big.matrix`, 20  
`write.big.matrix`, `big.matrix`, `character-method`  
    (`big.matrix-class`), 9